

#Question 1: Compare Hadoop and Spark:

Hadoop is designed to handle batch processing efficiently. Spark is designed to handle real-time data efficiently.

Hadoop is a high latency computing framework, which does not have an interactive mode.

Spark is a low latency computing and can process data interactively.

#Question 2: What is Apache Spark?

Apache Spark is an open-source, distributed processing system used for big data workloads. It utilizes in-memory caching, and optimized query execution for fast analytic queries against data of any size.

#Question 3: Explain the key features of Apache Spark.

Apache Spark is an open-source, distributed processing system used for big data workloads. It utilizes in-memory caching, and optimized query execution for fast analytic queries against data of any size.

#Question 4: What are the languages supported by Apache Spark and which is the most popular one?

Scala, Java, Python and R. Among these languages, Scala and Python have interactive shells for Spark.

The Scala shell can be accessed through spark-shell and the Python shell through pyspark.

Scala is the most used among them because Spark is written in Scala and it is the most popularly used for Spark.

#Question 5: What are benefits of Spark over MapReduce?

The primary difference between Spark and MapReduce is that Spark processes and retains data in memory for subsequent steps, whereas MapReduce processes data on disk.

As a result, for smaller workloads, Spark's data processing speeds are up to 100x faster than MapReduce.

#Question 6: Explain the concept of Resilient Distributed Dataset (RDD):

RDD was the primary user-facing API in Spark since its inception. At the core, an RDD is an immutable distributed collection of elements of your data, partitioned across nodes in your cluster that can be operated in parallel with a low-level API that offers transformations and actions.

#Question 7: How do we create RDDs in Spark?

There are two ways to create RDDs: parallelizing an existing collection in your driver program, or referencing a dataset in an external storage system, such as a shared filesystem, HDFS, HBase, or any data source offering a Hadoop InputFormat.

#Question 8: What is Executor Memory in a Spark application?

An executor is a process that is launched for a Spark application on a worker node. Each executor memory is the sum of yarn overhead memory and JVM Heap memory. JVM Heap memory comprises of: RDD Cache Memory. Shuffle Memory

#Question 9: What do you understand by Transformations in Spark?

Spark Transformation is a function that produces new RDD from the existing RDDs. It takes RDD as input and produces one or more RDD as output. Each time it creates new RDD when we apply any transformation. Thus, the so input RDDs, cannot be changed since RDD are immutable in nature.

#Question 10: Define Actions in Spark:

Actions are RDD's operation, that value returns back to the spark driver programs, which kick off a job to execute on a cluster.

Transformation's output is an input of Actions.

reduce, collect, takeSample, take, first, saveAsTextfile, saveAsSequenceFile, countByKey, foreach are common actions in Apache spark.

